

```
In [2]: import tensorflow as tf  
from tensorflow.keras import models, layers
```

here we will load the data into tensor flow

```
In [4]: print(tf.config.list_physical_devices('GPU'))  
[PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
In [7]: IMAGE_SIZE = 256  
BATCH_SIZE = 16  
CHANNELS = 3  
EPOCHS = 50
```

```
In [9]: dataset = tf.keras.preprocessing.image_dataset_from_directory(  
    "PlantVillage/Plant",  
    shuffle = True,  
    image_size = (IMAGE_SIZE,IMAGE_SIZE),  
    batch_size = BATCH_SIZE  
)
```

Found 13692 files belonging to 16 classes.

```
In [11]: def get_partition_and_data_tuning(ds,train_split= 0.8 ,val_split=0.1,test_split = 0  
train_size = int(len(dataset)*train_split)  
val_size = int(len(dataset)*val_split)  
test_size = int(len(dataset)*test_split)  
if shuffle:  
    ds = ds.shuffle(shuffle_size,seed = 12)  
train_data = ds.take(train_size)  
remain_data = ds.skip(train_size)  
val_data = remain_data.take(val_size)  
test_data = remain_data.skip(val_size)  
train_data = train_data.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)  
val_data = val_data.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)  
test_data = test_data.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)  
return train_data,val_data,test_data
```

```
In [13]: train, val, test = get_partition_and_data_tuning(dataset)
```

```
In [15]: resize_and_rescale = tf.keras.Sequential([  
  
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE,IMAGE_SIZE),  
    layers.experimental.preprocessing.Rescaling(1.0/255),  
])
```

```
In [17]: data_augmentation = tf.keras.Sequential([  
    layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),  
    layers.experimental.preprocessing.RandomRotation(0.2),  
])
```

```
In [19]: input=(IMAGE_SIZE,IMAGE_SIZE,CHANNELS)
n_classes = 11
```

```
In [21]: model = tf.keras.Sequential([
    layers.Conv2D(32,(3,3),activation="relu",input_shape=input),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,(3,3),activation="relu"),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,(3,3),activation="relu"),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,(3,3),activation="relu"),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,(3,3),activation="relu"),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64,(3,3),activation="relu"),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(64,activation="relu"),
    layers.Dense(n_classes,activation="softmax")
])
```

```
In [23]: model.build(input_shape = input)
```

```
In [25]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_3 (Conv2D)	(None, 28, 28, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_4 (Conv2D)	(None, 12, 12, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_5 (Conv2D)	(None, 4, 4, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 64)	16448
dense_1 (Dense)	(None, 11)	715
=====		
Total params: 184,267		
Trainable params: 184,267		
Non-trainable params: 0		

```
In [ ]: with tf.device('/device:GPU:0'):  
    model.compile(  
        optimizer='adam',  
        loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),  
        metrics=['accuracy'])  
    history = model.fit(  
        train,  
        epochs = EPOCHS,  
        batch_size = BATCH_SIZE,
```

```
        verbose = 1,
        validation_data = val
    )

Epoch 1/50
271/684 [=====>.....] - ETA: 33s - loss: nan - accuracy: 0.0473

In [15]: scores = model.evaluate(test)

61/61 [=====] - 25s 53ms/step - loss: 0.1069 - accuracy: 0.9688

In [16]: scores

Out[16]: [0.10689208656549454, 0.9688149690628052]

In [17]: history.params

Out[17]: {'verbose': 1, 'epochs': 50, 'steps': 482}

In [18]: history.history.keys()

Out[18]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

In [19]: len(history.history['accuracy'])

Out[19]: 50

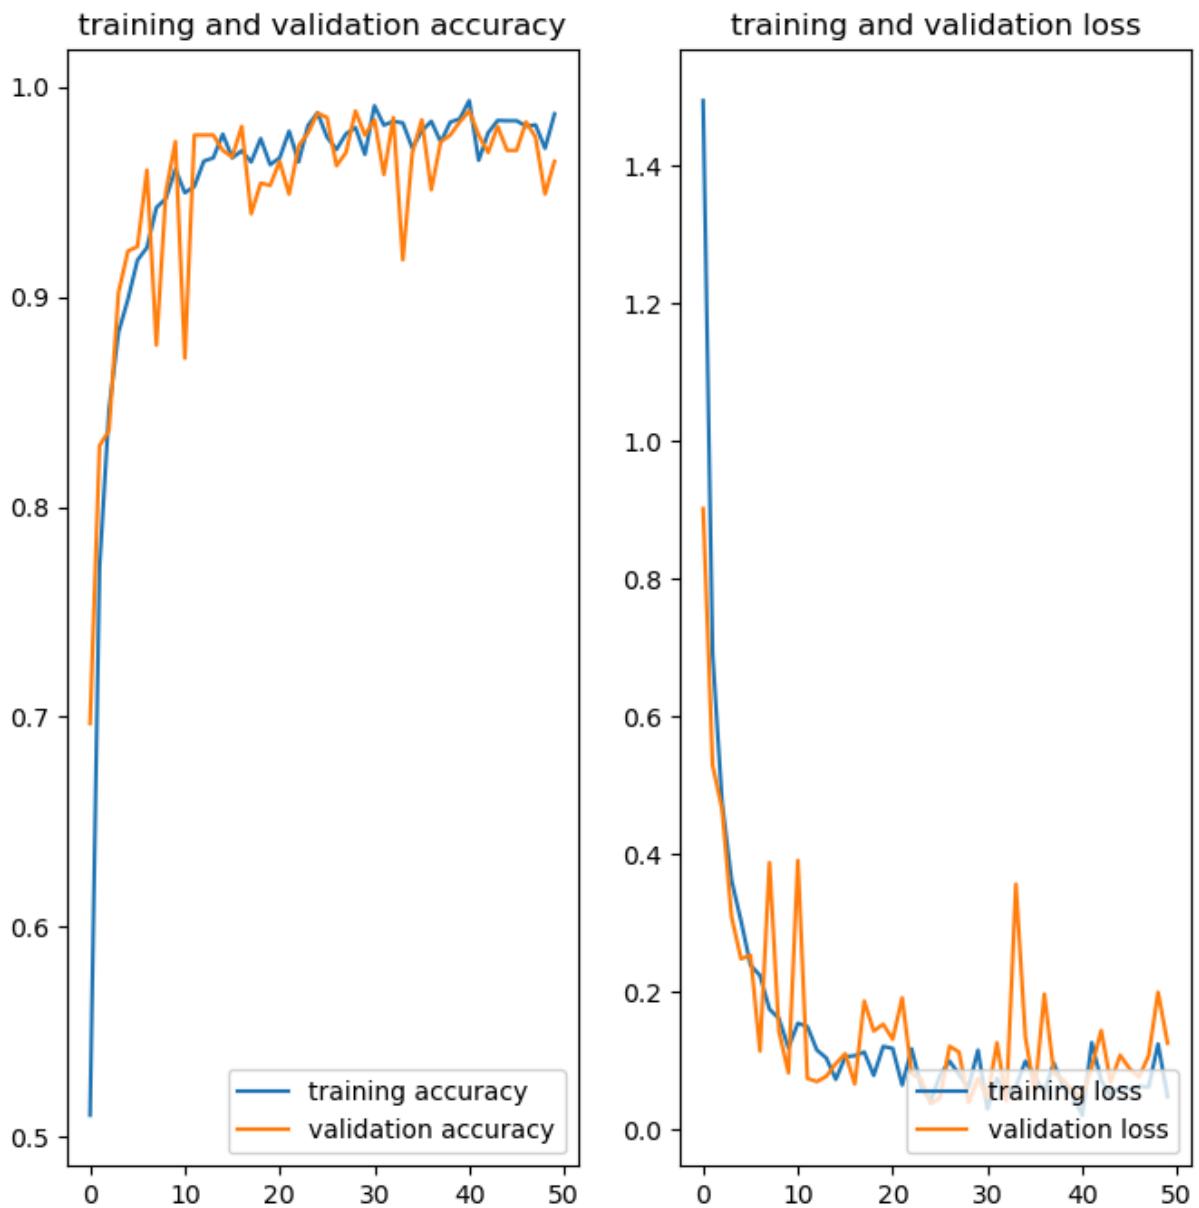
In [20]: import matplotlib as plt

In [21]: acc = history.history['accuracy']
        val_acc = history.history['val_accuracy']
        loss = history.history['loss']
        val_loss = history.history['val_loss']

In [22]: plt.figure(figsize=(8,8))
        plt.subplot(1,2,1)
        plt.plot(range(EPOCHS),acc,label="training accuracy")
        plt.plot(range(EPOCHS),val_acc,label="validation accuracy")
        plt.legend(loc="lower right")
        plt.title("training and validation accuracy")

        plt.subplot(1,2,2)
        plt.plot(range(EPOCHS),loss,label="training loss")
        plt.plot(range(EPOCHS),val_loss,label="validation loss")
        plt.legend(loc="lower right")
        plt.title("training and validation loss")

Out[22]: Text(0.5, 1.0, 'training and validation loss')
```



```
In [23]: import os
os.listdir("../leaf disease prediction/Models")
```

```
Out[23]: ['.ipynb_checkpoints',
 '2',
 'apple',
 'apple_aware_model',
 'cherry',
 'cherry_aware_model',
 'peach',
 'peach_aware_model',
 'pepper',
 'pepper_aware_model',
 'potato',
 'potato_aware_model',
 'quantize_aware_model',
 'straberry',
 'straberry_aware_model']
```

```
In [24]: # version = max([int(i) for i in os.listdir("../leaf disease prediction/Models")]+[  
model.save(f"../leaf disease prediction/Models/quantize")]
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_com  
piled_convolution_op, _jit_compiled_convolution_op, _jit_compiled_convolution_op, _j  
it_compiled_convolution_op while saving (showing 5 of 6). These functions will not b  
e directly callable after loading.  
INFO:tensorflow:Assets written to: ../leaf disease prediction/Models/quantize\assets  
INFO:tensorflow:Assets written to: ../leaf disease prediction/Models/quantize\assets
```

```
In [ ]:
```

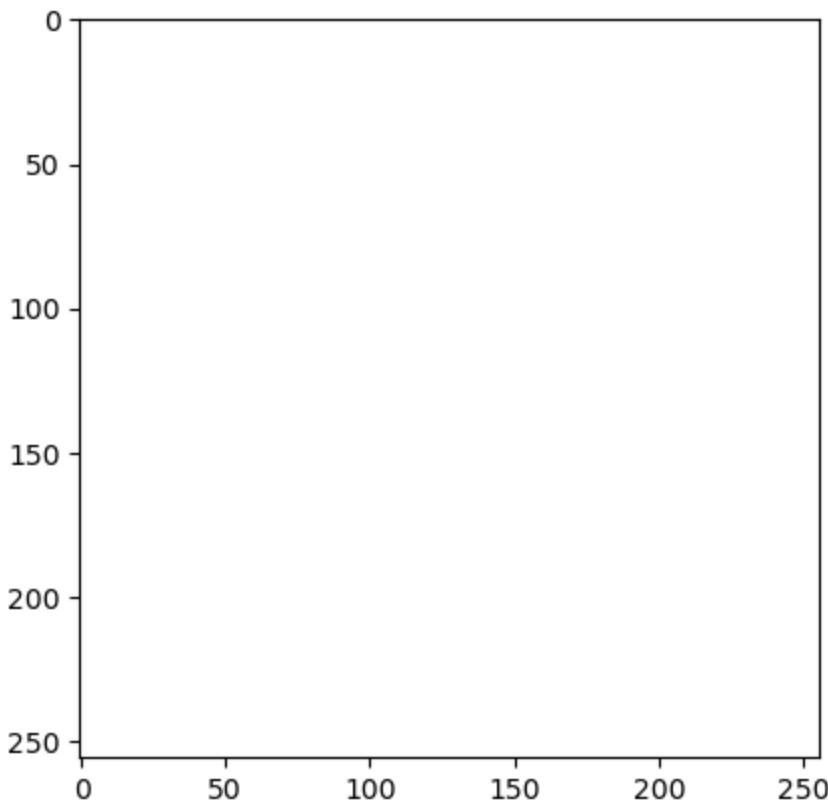
```
In [25]: # _____ predicting _____
```

```
In [26]: import numpy as np
```

```
In [27]: for images_batch,labels_batch in test.take(1):  
    first_img = images_batch[0].numpy().astype('float32')  
    first_label = labels_batch[0].numpy()  
    print("first image to predict")  
    plt.imshow(first_img)  
    print("actual label:",dataset.class_names[first_label])  
    batch_prediction = model.predict(images_batch)  
    label1 = np.argmax(batch_prediction[0])  
    print("predicted label:",dataset.class_names[label1])
```

```
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB  
data ([0..1] for floats or [0..255] for integers).
```

```
first image to predict  
actual label: Apple_Cedar_apple_rust  
1/1 [=====] - 0s 301ms/step  
predicted label: Apple_Cedar_apple_rust
```



```
In [28]: def predict(model,img):  
  
    img_array = tf.keras.preprocessing.image.img_to_array(img)  
    img_array = tf.expand_dims(img_array,0)#create a batch  
    prediction = model.predict(img_array)  
    pre_class = dataset.class_names[np.argmax(prediction[0])]  
    confidence = round(100*(np.max(prediction[0])),2)  
    return pre_class,confidence
```

```
In [29]: # _____Loding and using the saved model_____
```

```
In [30]: from tensorflow.keras.models import load_model  
model = load_model("../leaf disease prediction/Models/quantize")
```

```
In [ ]:
```

```
In [31]: import tensorflow_model_optimization as tfmot
```

```
In [32]: new_model = tf.keras.models.load_model('Models/quantize')
```

```
In [33]: quantize_model = tfmot.quantization.keras.quantize_model
```

```
In [34]: quantize_aware_model = quantize_model(new_model)
```

```
In [35]: quantize_aware_model.compile(  
    optimizer='adam',  
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),  
    metrics=['accuracy'])
```

```
In [36]: quantize_aware_model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
quantize_layer (Quantizer)	(None, 256, 256, 3)	3
quant_conv2d (QuantizeWrapperV2)	(None, 254, 254, 32)	963
quant_max_pooling2d (QuantizeWrapperV2)	(None, 127, 127, 32)	1
quant_conv2d_1 (QuantizeWrapperV2)	(None, 125, 125, 64)	18627
quant_max_pooling2d_1 (QuantizeWrapperV2)	(None, 62, 62, 64)	1
quant_conv2d_2 (QuantizeWrapperV2)	(None, 60, 60, 64)	37059
quant_max_pooling2d_2 (QuantizeWrapperV2)	(None, 30, 30, 64)	1
quant_conv2d_3 (QuantizeWrapperV2)	(None, 28, 28, 64)	37059
quant_max_pooling2d_3 (QuantizeWrapperV2)	(None, 14, 14, 64)	1
quant_conv2d_4 (QuantizeWrapperV2)	(None, 12, 12, 64)	37059
quant_max_pooling2d_4 (QuantizeWrapperV2)	(None, 6, 6, 64)	1
quant_conv2d_5 (QuantizeWrapperV2)	(None, 4, 4, 64)	37059
quant_max_pooling2d_5 (QuantizeWrapperV2)	(None, 2, 2, 64)	1
quant_flatten (QuantizeWrapperV2)	(None, 256)	1
quant_dense (QuantizeWrapperV2)	(None, 64)	16453
quant_dense_1 (QuantizeWrapperV2)	(None, 11)	720
=====		
Total params: 185,009		
Trainable params: 184,267		

```
Non-trainable params: 742
```

```
In [37]: with tf.device('/device:GPU:0'):  
    quantize_aware_model.fit(  
        train,  
        epochs = 10,  
        batch_size = BATCH_SIZE,  
        verbose = 1,  
        validation_data = val  
    )
```

```
Epoch 1/10  
482/482 [=====] - 108s 219ms/step - loss: 0.2550 - accuracy: 0.9224 - val_loss: 0.0998 - val_accuracy: 0.9719  
Epoch 2/10  
482/482 [=====] - 77s 159ms/step - loss: 0.0975 - accuracy: 0.9729 - val_loss: 0.2510 - val_accuracy: 0.9375  
Epoch 3/10  
482/482 [=====] - 69s 143ms/step - loss: 0.1286 - accuracy: 0.9725 - val_loss: 0.2399 - val_accuracy: 0.9521  
Epoch 4/10  
482/482 [=====] - 75s 154ms/step - loss: 0.0825 - accuracy: 0.9808 - val_loss: 0.0680 - val_accuracy: 0.9854  
Epoch 5/10  
482/482 [=====] - 68s 142ms/step - loss: 0.0583 - accuracy: 0.9834 - val_loss: 0.1412 - val_accuracy: 0.9688  
Epoch 6/10  
482/482 [=====] - 74s 154ms/step - loss: 0.0877 - accuracy: 0.9799 - val_loss: 0.1356 - val_accuracy: 0.9667  
Epoch 7/10  
482/482 [=====] - 79s 163ms/step - loss: 0.0495 - accuracy: 0.9856 - val_loss: 0.1833 - val_accuracy: 0.9490  
Epoch 8/10  
482/482 [=====] - 73s 151ms/step - loss: 0.0581 - accuracy: 0.9852 - val_loss: 0.1010 - val_accuracy: 0.9802  
Epoch 9/10  
482/482 [=====] - 74s 154ms/step - loss: 0.0562 - accuracy: 0.9845 - val_loss: 0.0740 - val_accuracy: 0.9802  
Epoch 10/10  
482/482 [=====] - 65s 134ms/step - loss: 0.0235 - accuracy: 0.9923 - val_loss: 0.0704 - val_accuracy: 0.9844
```

```
In [38]: ##### quantize_aware_model.evaluate(test)  
quantize_aware_model.save(f"../leaf disease prediction/Models/quantize_aware_
```

```
WARNING:absl:Found untraced functions such as conv2d_layer_call_fn, conv2d_layer_call_and_return_conditional_losses, _jit_compiled_convolution_op, conv2d_1_layer_call_fn, conv2d_1_layer_call_and_return_conditional_losses while saving (showing 5 of 24). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: ../leaf disease prediction/Models/quantize_aware_model\assets
```

```
INFO:tensorflow:Assets written to: ../leaf disease prediction/Models/quantize_aware_model\assets
```

```
In [39]: convertor = tf.lite.TFLiteConverter.from_keras_model(quantize_aware_model)
convertor.optimizations = [tf.lite.Optimize.DEFAULT]
tf_model1 = convertor.convert()
```

```
WARNING:absl:Found untraced functions such as conv2d_layer_call_fn, conv2d_layer_call_and_return_conditional_losses, _jit_compiled_convolution_op, conv2d_1_layer_call_fn, conv2d_1_layer_call_and_return_conditional_losses while saving (showing 5 of 24). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: C:\Users\heman\AppData\Local\Temp\tmpbpmpwed5\assets
```

```
INFO:tensorflow:Assets written to: C:\Users\heman\AppData\Local\Temp\tmpbpmpwed5\assets
```

```
C:\Users\heman\anaconda3\lib\site-packages\tensorflow\lite\python\convert.py:766: UserWarning: Statistics for quantized inputs were expected, but not specified; continuing anyway.
```

```
    warnings.warn("Statistics for quantized inputs were expected, but not "
```

```
In [40]: with open("quantize_tflite_model.tflite","wb") as f:
    f.write(tf_model1)
```

```
In [3]: nbconvert --to webpdf --allow-chromium-download modeltraining.ipynb
```

```
Cell In[3], line 1
```

```
nbconvert --to webpdf --allow-chromium-download modeltraining.ipynb
```

```
^
```

```
SyntaxError: invalid syntax
```

```
In [ ]:
```